

Normal Two-Way ANOVA

OUTLINE

10.1 Introduction: Main and Interaction Effects	129
10.2 Data Generation	131
10.3 Aside: Using Simulation to Assess Bias and Precision of an Estimator	133
10.4 Analysis Using R	134
10.5 Analysis Using WinBUGS	135
10.5.1 Main-Effects ANOVA Using WinBUGS	135
10.5.2 Interaction-Effects ANOVA Using WinBUGS	137
10.5.3 Forming Predictions	138
10.6 Summary	139

10.1 INTRODUCTION: MAIN AND INTERACTION EFFECTS

We now extend the one-way analysis of variance (ANOVA) model by adding another factor and arrive at a two-way ANOVA. We only consider fixed effects here. There are two ways in which the effects of two factors A and B can be combined, and the associated models are called main-effects and interaction-effects model. In the main-effects model, the effects of A and B are additive, i.e., the effect of one level of factor A, say a_1 , does not depend on whether it is assessed at one level of B, say b_1 , or at another, say b_2 . In contrast, with an interaction between A and B, some or all effects depend on some or all of each other and the effect of a_1 may not be identical when assessed at b_1 or at b_2 . Interaction is symmetric, so the effect of b_1 will in general also not be the same whether assessed at a_1 or at a_2 . However, the interaction model is still linear, since effects are simply

added together, only with an additional set of effects: those for the combination of each level of two or more factors. When factors are considered fixed, then depending on the model, not all effects will be estimable, see later (10.5.1). In contrast, in a random-effects model with interaction, all effects will in general be estimable.

In Chapter 6 we already saw the linear models for the two-way ANOVA with interaction. In short, the effects parameterization for a model with two factors A (with j levels) and B (with k levels) is this:

$$y_i = \alpha + \beta_{j(i)} * A_i + \delta_{k(i)} * B_i + \gamma_{jk(i)} * A_i * B_i + \varepsilon_i,$$

while the means parameterization is this:

$$y_i = \alpha_{jk(i)} * A_i * B_i + \varepsilon_i$$

In both cases, we need to assume a distribution for the residuals to complete the model description:

$$\varepsilon_i \sim \text{Normal}(0, \sigma^2).$$

We will use the beautiful mourning cloak (Fig. 10.1) as an illustration for this chapter and assume that we measured wing length of butterflies in each of three elevation classes in five different populations and that the effects of these factors interact. Table 10.1 shows the meaning of the coefficients in the linear model for the effects parameterization.



FIGURE 10.1 Mourning cloak (*Nymphalis antiopa*), Switzerland, 2006. (Photo: T. Marent)

TABLE 10.1 The 15 Parameters Estimated in the Effects Parameterization of a Two-Way ANOVA with Interaction for the Mourning Cloak Example.

Intercept	Elevation2	Elevation3
pop2	pop2.elevation2	pop2.elevation3
pop3	pop3.elevation2	pop3.elevation3
pop4	pop4.elevation2	pop4.elevation3
pop5	pop5.elevation2	pop5.elevation3

If the population factor has $n.\text{pop}$ levels and the elevation factor $n.\text{elev}$ levels, we have one intercept, $n.\text{pop}-1 = 4$ effects for the population factor, $n.\text{elev}-1 = 2$ effects for the elevation factor and $(n.\text{pop}-1) * (n.\text{elev}-1) = 8$ effects for the interaction between population and elevation. This adds up to the 15 degrees of freedom that it will cost us to fit this model to a data set that contains observations in every cell (combination of levels) in the cross-classification of population and elevation.

10.2 DATA GENERATION

We assume five populations with 12 butterflies were measured in each and that of these 12, four butterflies were studied at each of three elevation classes (low, medium, high). Wing length differs with elevation, perhaps because butterflies hatch at different size at different elevation or because of different size-dependent predation owing to different bird communities at different elevations. Furthermore, the relationship between wing length and elevation class is not homogeneous among the five studied populations so there is a population-elevation interaction. Residual wing length standard deviation will be 3.

```
# Choose sample size
n.pop <- 5
n.elev <- 3
nsample <- 12
n <- n.pop * nsample

# Create factor levels
pop <- gl(n = n.pop, k = nsample, length = n)
elev <- gl(n = n.elev, k = nsample / n.elev, length = n)

# Choose effects
baseline <- 40                # Intercept
pop.effects <- c(-10, -5, 5, 10) # Population effects
elev.effects <- c(5, 10)       # Elev effects
```




FIGURE 10.2 Mean wing length of mourning cloaks at each of three elevations and in each of five populations. Boxplots are ordered first by elevation and second by population and their identity is recognizable from the tick labels on the x-axis. For instance, the boxplot labelled 3.2 on the x-axis shows the mean wing length in population 2 at elevation 3.

```
interaction.effects <- c(-2, 3, 0, 4, 4, 0, 3, -2) # Interaction effects
all.effects <- c(baseline, pop.effects, elev.effects, interaction.effects)

sigma <- 3
eps <- rnorm(n, 0, sigma) # Residuals

X <- as.matrix(model.matrix(~ pop*elev) ) # Create design matrix
X # Have a look at that
```

Use matrix multiplication to assemble all components for the final wing length measurements y which we inspect in a grouped boxplot (Fig. 10.2).

```
wing <- as.numeric(as.matrix(X) %*% as.matrix(all.effects) + eps)
# NOTE: as.numeric is ESSENTIAL for WinBUGS later
boxplot(wing ~ elev*pop, col = "grey", xlab = "Elevation-by-Population", ylab =
"Wing length", main = "Simulated data set", las = 1, ylim = c(20, 70)) # Plot of
generated data
abline(h = 40)
```

We have generated data for which the wing length–elevation relationship varies considerably among the five populations. This can also be nicely seen in a useful conditioning plot, which can be drawn using the function `xypplot()` in the `lattice` package. The data can be viewed in two ways; both plots show that the effects of population and elevation are not independent.

```
library("lattice")           # Load the lattice library
xyplot(wing ~ elev | pop, ylab = "Wing length", xlab = "Elevation", main =
"Population-specific relationship between wing and elevation class")
xyplot(wing ~ pop | elev, ylab = "Wing length", xlab = "Population", main =
"Elevation-specific relationship between wing and population")
```

10.3 ASIDE: USING SIMULATION TO ASSESS BIAS AND PRECISION OF AN ESTIMATOR

Let's quickly compare our parameter estimates with what we put into the data:

```
lm(wing ~ pop*elev)
all.effects

> lm(wing ~ pop*elev)

Call:
lm(formula = wing ~ pop * elev)

Coefficients:
(Intercept)      pop2      pop3      pop4      pop5      elev2
    38.859    -8.543    -4.793     4.702    13.410     6.437
      elev3 pop2:elev2 pop3:elev2 pop4:elev2 pop5:elev2 pop2:elev3
    11.213    -1.284     1.530     0.332     1.857     1.359
 pop3:elev3 pop4:elev3 pop5:elev3
     1.820     2.835    -6.609

> all.effects
[1] 40 -10 -5  5 10  5 10 -2  3  0  4  4  0  3 -2
```

The coefficient estimates don't necessarily resemble very much the parameters from which we simulated these data; after all, our sample size is rather small. So, to reassure ourselves that these differences are simply due to sampling variation, we repeat this data generation-analysis cycle 1000 times and average over the random sampling variation to convince ourselves that the estimators from the linear model are indeed unbiased. Simulations of this kind can be done easily in R, and this is one of the great strengths of R.

```
n.iter <- 1000           # Desired number of iterations
estimates <- array(dim = c(n.iter, length(all.effects))) # Data structure to
hold results

for(i in 1:n.iter) {     # Run simulation n.iter times
  print(i)               # Optional
  eps <- rnorm(n, 0, sigma) # Residuals
```

```

y <- as.numeric(as.matrix(X) %*% as.matrix(all.effects) + eps) # Assemble data
fit.model <- lm(y ~ pop*elev)                                # Break down data
estimates[i,] <- fit.model$coefficients                     # Save estimates of coefs.
}

```

Compare the input (i.e., the chosen effects) and the output when averaged over sampling variation:

```

print(apply(estimates, 2, mean), dig = 2)
all.effects

> print(apply(estimates, 2, mean), dig = 2)
[1] 40.0102 -10.0452 -5.0717 4.9687 9.9524 5.0417 10.0419 -1.9761
[9] 2.9493 -0.0226 4.0031 3.9407 -0.0039 3.0234 -1.9838
> all.effects
[1] 40 -10 -5 5 10 5 10 -2 3 0 4 4 0 3 -2

```

These are much closer to our input parameters. Depending on the number of iterations, we can get arbitrarily close to the input. Alternatively, we could increase the sample size from 12 to 12000, and we would get estimates that are still closer to the input values.

10.4 ANALYSIS USING R

We continue with the analysis of our data set and use R to fit the main-effects model first.

```

mainfit <- lm(wing ~ elev + pop)
mainfit

> mainfit
[ ... ]
Coefficients:
(Intercept)      elev2      elev3      pop2
      38.736      6.924      11.095     -8.518
      pop3      pop4      pop5
     -3.676      5.757      11.826

```

Then, we fit the means parameterization of the interaction model.

```

intfit <- lm(wing ~ elev*pop-1-pop-elev)
intfit

> intfit <- lm(wing ~ elev*pop-1-pop-elev)
> intfit
[ ... ]

```

Coefficients:

elev1:pop1	elev2:pop1	elev3:pop1
38.86	45.30	50.07
elev1:pop2	elev2:pop2	elev3:pop2
30.32	35.47	42.89
elev1:pop3	elev2:pop3	elev3:pop3
34.07	42.03	47.10
elev1:pop4	elev2:pop4	elev3:pop4
43.56	50.33	57.61
elev1:pop5	elev2:pop5	elev3:pop5
52.27	60.56	56.87

10.5 ANALYSIS USING WinBUGS

10.5.1 Main-Effects ANOVA Using WinBUGS

We fit the main-effects model in the effects parameterization because I find that easier to code. One minor feature in this analysis is the way in which we specify the priors for the elements of the parameter vectors: instead of looping over each of them, we now write them all out, since we have to set to zero the first (or another) level of each factor to make this fixed-effects model identifiable.

```
# Define model
sink("2w.anova.txt")
cat("
model {

# Priors
  alpha ~ dnorm(0, 0.001)           # Intercept
  beta.pop[1] <- 0                  # set to zero effect of 1st level
  beta.pop[2] ~ dnorm(0, 0.001)
  beta.pop[3] ~ dnorm(0, 0.001)
  beta.pop[4] ~ dnorm(0, 0.001)
  beta.pop[5] ~ dnorm(0, 0.001)
  beta.elev[1] <- 0                  # ditto
  beta.elev[2] ~ dnorm(0, 0.001)
  beta.elev[3] ~ dnorm(0, 0.001)
  sigma ~ dunif(0, 100)

# Likelihood
  for (i in 1:n) {
    wing[i] ~ dnorm(mean[i], tau)
    mean[i] <- alpha + beta.pop[pop[i]] + beta.elev[elev[i]]
  }
}
```

Derived quantities

```
tau <- 1 / ( sigma * sigma )
}
",fill=TRUE)
sink()
```

Bundle data

```
win.data <- list(wing=wing, elev = as.numeric(elev), pop = as.numeric(pop), n =
length(wing))
```

Inits function

```
inits <- function(){ list(alpha = rnorm(1), sigma = rlnorm(1) ) }
```

Parameters to estimate

```
params <- c("alpha", "beta.pop", "beta.elev", "sigma")
```

MCMC settings

```
ni <- 1200
nb <- 200
nt <- 2
nc <- 3
```

Start Gibbs sampling

```
out <- bugs(win.data, inits, params, "2w.anova.txt", n.thin=nt, n.chains=nc,
n.burnin=nb, n.iter=ni, debug = TRUE)
```

Print estimates

```
print(out, dig = 3)
> print(out, dig = 3)
```

```
Inference for Bugs model at "2w.anova.txt", fit using WinBUGS,
3 chains, each with 1200 iterations (first 200 discarded), n.thin = 2
n.sims = 1500 iterations saved
```

	mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
alpha	38.701	1.216	36.355	37.870	38.760	39.510	40.945	1.006	680
beta.pop[2]	-8.447	1.476	-11.325	-9.492	-8.442	-7.428	-5.547	1.001	1500
beta.pop[3]	-3.567	1.458	-6.360	-4.512	-3.586	-2.664	-0.670	1.001	1500
beta.pop[4]	5.820	1.470	3.040	4.825	5.871	6.788	8.713	1.003	780
beta.pop[5]	11.841	1.454	9.090	10.850	11.860	12.770	14.725	1.001	1500
beta.elev[2]	6.901	1.156	4.714	6.106	6.923	7.671	9.184	1.004	560
beta.elev[3]	11.077	1.119	8.988	10.300	11.040	11.850	13.220	1.002	970
sigma	3.561	0.360	2.954	3.306	3.538	3.788	4.323	1.000	1500
deviance	320.983	4.308	314.900	317.900	320.400	323.200	331.600	1.001	1500
[...]									

```
DIC info (using the rule, pD = Dbar-Dhat)
```

```
pD = 7.8 and DIC = 328.8
```

```
DIC is an estimate of expected predictive error (lower deviance is better).
```


We get estimates that are fairly similar with the MLEs above. To see the estimate of the residual, you can type `summary(mainfit)`.

10.5.2 Interaction-Effects ANOVA Using WinBUGS

We will specify the means parameterization for ease of coding and show how parameters in WinBUGS can be arrays with two (or more) dimensions. This is handy when organizing an analysis.

```
# Write model
sink("2w2.anova.txt")
cat("
model {

# Priors
  for (i in 1:n.pop){
    for(j in 1:n.elev) {
      group.mean[i,j] ~ dnorm(0, 0.0001)
    }
  }
  sigma ~ dunif(0, 100)

# Likelihood
  for (i in 1:n) {
    wing[i] ~ dnorm(mean[i], tau)
    mean[i] <- group.mean[pop[i], elev[i]]
  }

# Derived quantities
  tau <- 1 / ( sigma * sigma)
}
",fill=TRUE)
sink()

# Bundle data
win.data <- list(wing=wing, elev = as.numeric(elev), pop = as.numeric(pop), n =
length(wing), n.elev = length(unique(elev)), n.pop = length(unique(pop)))

# Inits function
inits <- function(){list(sigma = rlnorm(1) )}

# Parameters to estimate
params <- c("group.mean", "sigma")

# MCMC settings
ni <- 1200
nb <- 200
nt <- 2
nc <- 3
```

```
# Start Gibbs sampling
out <- bugs(win.data, inits, params, "2w2.anova.txt", n.thin=nt, n.chains=nc,
n.burnin=nb, n.iter=ni, debug = TRUE)

# Print estimates
print(out, dig = 3)
> print(out, dig = 3)

Inference for Bugs model at "2w2.anova.txt", fit using WinBUGS,
  3 chains, each with 1200 iterations (first 200 discarded), n.thin = 2
n.sims = 1500 iterations saved
```

	mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
group.mean[1,1]	38.864	1.644	35.500	37.790	38.885	39.982	42.040	1.001	1500
group.mean[1,2]	45.263	1.600	42.135	44.187	45.250	46.300	48.506	1.000	1500
group.mean[1,3]	50.062	1.645	46.784	48.987	50.080	51.110	53.330	1.000	1500
group.mean[2,1]	30.330	1.589	27.150	29.297	30.320	31.370	33.495	1.000	1500
group.mean[2,2]	35.405	1.637	32.284	34.340	35.390	36.532	38.495	1.004	720
group.mean[2,3]	42.905	1.564	39.708	41.865	42.930	43.950	45.850	1.003	780
group.mean[3,1]	34.071	1.634	30.975	32.977	34.065	35.120	37.346	1.002	1500
group.mean[3,2]	41.932	1.609	38.714	40.897	41.940	43.010	45.041	1.001	1500
group.mean[3,3]	47.077	1.628	43.855	46.017	47.040	48.160	50.290	1.004	500
group.mean[4,1]	43.508	1.663	40.345	42.370	43.440	44.580	46.895	1.001	1500
group.mean[4,2]	50.403	1.635	47.239	49.310	50.365	51.520	53.590	1.004	540
group.mean[4,3]	57.630	1.617	54.525	56.570	57.640	58.720	60.770	1.002	1500
group.mean[5,1]	52.246	1.608	49.054	51.177	52.260	53.350	55.411	1.002	1100
group.mean[5,2]	60.538	1.606	57.384	59.457	60.580	61.630	63.625	1.000	1500
group.mean[5,3]	56.893	1.621	53.670	55.877	56.870	57.962	60.095	1.002	1300
sigma	3.227	0.358	2.617	2.980	3.193	3.447	4.014	1.001	1500
deviance	309.326	6.734	299.100	304.500	308.350	313.000	324.952	1.001	1500
[...]									

```
DIC info (using the rule, pD = Dbar-Dhat)
pD = 15.7 and DIC = 325.1
DIC is an estimate of expected predictive error (lower deviance is better).
```

We find the usual similarity between the Bayes and the maximum likelihood solution above (do `summary(intfit)`) and note in passing that the estimated number of parameters (pD) is pretty close to what we would expect it to be.

10.5.3 Forming Predictions

Let's present the Bayesian inference for the interaction-effects model in a graph showing the predicted response, analogous to least-square means in a classical analysis, for each combination of elevation and population (Fig. 10.3). This plot corresponds to the boxplot of the data set (Fig. 10.2); or selects the order of the predictions to match that in Fig. 10.2.




FIGURE 10.3 Predicted wing length of mourning cloaks for each elevation-population combination. Error bars are 1 SE.

```
or <- c(1,4,7,10,13,2,5,8,11,14,3,6,9,12,15)
plot(or, out$mean$group.mean, xlab = "Elev-by-Population", las = 1, ylab =
"Predicted wing-length", cex = 1.5, ylim = c(20, 70))
segments(or, out$mean$group.mean, or, out$mean$group.mean + out$sd$group.mean,
col = "black", lwd = 1)
segments(or, out$mean$group.mean, or, out$mean$group.mean - out$sd$group.mean,
col = "black", lwd = 1)
abline(h = 40)
```

10.6 SUMMARY

We have introduced the concepts of main and interaction effects and used R and WinBUGS to fit the corresponding two-way ANOVA models. In an aside, we have illustrated R's flexibility to conduct simulations to verify the effects of sampling variation on the parameter estimates.

EXERCISES

1. *Toy snake example*: Fit a two-way ANOVA with interaction to the toy example of Chapter 6 and see what happens to the nonidentifiable parameter.
2. *Swiss hare data*: Fit an ANOVA model to mean hare density to decide whether the effect of grassland and arable land use is the same in all regions. Regions and land use are somewhat confounded, but we ignore this here.