

General Linear Model (ANCOVA)

OUTLINE

11.1 Introduction	141
11.2 Data Generation	143
11.3 Analysis Using R	145
11.4 Analysis Using WinBUGS (And A Cautionary Tale About the Importance of Covariate Standardization)	145
11.5 Summary	149

11.1 INTRODUCTION

The “model of the mean,” t-test, simple linear regression, and analysis of variance (ANOVA) are all just special cases of a very general and powerful statistical model, the general linear model (\neq GLM!; see Chapter 13). This model expresses a continuous response as a linear combination of the effects of discrete and/or continuous explanatory variables plus a single random contribution from a normal distribution, whose variance is estimated along with the coefficients of all discrete and continuous covariates and possible interactions.

Before this was widely recognized, people used to make a rather sharp and artificial distinction between linear models that contain categorical explanatory variables only and were called t-test or ANOVA models and those that contain continuous covariates only and were called regression models. Models that contained both types of explanatory variables were usually treated as ANOVAs with typically a single continuous covariate to correct for preexisting variation among experimental units. These



FIGURE 11.1 Male asp viper (*Vipera aspis*), Switzerland, 2007. (Photo T. Ott)

models were called analysis of covariance (ANCOVA) models and that's why I use this term here.

Nowadays, in many practical applications, we typically have several explanatory variables of both types. In addition, we want to fit both main effects of these covariates and some or all their pairwise or even higher-order interactions. As an example, we here consider an ANCOVA model with an interaction between a discrete and a continuous covariate. We saw how to fit interactions between two discrete covariates in Chapter 10. Interactions between two continuous covariates are easy to fit: simply fit one additional covariate whose values are obtained by multiplication of the two main covariates.

In this chapter, we consider the relationship between body mass and body length of the asp viper (Fig. 11.1) in three populations: Pyrenees, Massif Central, and the Jura mountains. We are particularly interested in population-specific differences of the mass–length relationship, i.e., in the interaction between length and population. The means parameterization of the model we will fit can be written as (see Section 6.3.6)

$$y_i = \alpha_{j(i)} + \beta_{j(i)} * x_i + \varepsilon_i$$

and

$$\varepsilon_i \sim \text{Normal}(0, \sigma^2),$$

where y_i is the body mass of individual i , $\alpha_{j(i)}$ and $\beta_{j(i)}$ are the intercept and the slope, respectively, of the mass–length relationship in population j , x_i is the body length of snake i , and as usual, ε_i describes the combined effects of all unmeasured influences on the body mass of snake i and is assumed to behave like a normal random variable whose variance σ^2 we estimate.

The effects parameterization of the same model is this:

$$y_i = \alpha_{\text{Pyr}} + \beta_1 * x_{\text{MC}(i)} + \beta_2 * x_{\text{Jura}(i)} + \beta_3 * x_{\text{body}(i)} + \beta_4 * x_{\text{body}(i)} * x_{\text{MC}(i)} \\ + \beta_5 * x_{\text{body}(i)} * x_{\text{Jura}(i)} + \varepsilon_i$$

In addition to y_i and ε_i that are as before, α_{Pyr} is the expected mass of snakes in the Pyrenees, β_1 is the difference between the expected mass of snakes in the Massif Central and that in the Pyrenees, and $x_{\text{MC}(i)}$ is the indicator for snakes caught in the Massif Central. β_2 is the difference between the expected mass in the Jura and that in the Pyrenees, $x_{\text{Jura}(i)}$ is the indicator for snakes in the Jura, β_3 is the slope of the regression of body mass on body length x_{body} in the Pyrenees, β_4 is the difference in that slope between the Massif Central and the Pyrenees, and β_5 is the difference of slopes between the Jura and the Pyrenees. Thus, snakes in the Pyrenees act as baseline with which snakes from the Massif Central and the Jura are compared, but as usual, this choice has no effect on inference.

11.2 DATA GENERATION

As always, we assume a balanced design simply for convenience of data generation.

```
n.groups <- 3
n.sample <- 10
n <- n.groups * n.sample # Total number of data points
x <- rep(1:n.groups, rep(n.sample, n.groups)) # Indicator for
population
pop <- factor(x, labels = c("Pyrenees", "Massif Central", "Jura"))
length <- runif(n, 45, 70) # Obs. body length (cm) is rarely less
than 45
```

We build the design matrix of an interactive combination of length and population, inspect that, and select the parameter values, i.e., choose values for α_{Pyr} , β_1 , β_2 , β_3 , β_4 , and β_5 .

```
Xmat <- model.matrix(~ pop*length)
print(Xmat, dig = 2)
beta.vec <- c(-250, 150, 200, 6, -3, -4)
```

Next, we build up the body mass measurements y_i by adding the residual to the value of the linear predictor, with residuals drawn from an appropriate zero-mean normal distribution. The value of the linear predictor is obtained by matrix multiplication of the design matrix (`Xmat`) and the parameter vector (`beta.vec`). Our vipers are probably all too fat, but that doesn't really matter for our purposes.

```
lin.pred <- Xmat[,] %*% beta.vec      # Value of lin.predictor
eps <- rnorm(n = n, mean = 0, sd = 10) # residuals
mass <- lin.pred + eps               # response = lin.pred + residual
hist(mass)                          # Inspect what we've created
matplot(cbind(length[1:10], length[11:20], length[21:30]), cbind(mass[1:10],
mass[11:20], mass[21:30]), ylim = c(0, max(mass)), ylab = "Body mass (g)", xlab =
"Body length (cm)", col = c("Red", "Green", "Blue"), pch = c("P", "M", "J"), las = 1,
cex = 1.2, cex.lab = 1.5)
```

We have created a data set in which vipers from the Pyrenees have the steepest slope between mass and length, followed by those from the Massif Central and finally those in the Jura mountains (Fig. 11.2). Now let's disassemble these data.




FIGURE 11.2 Simulated data set showing body mass versus length of 10 asp vipers in each of three populations (P – Pyrenees, M – Massif Central, and J – Jura).

11.3 ANALYSIS USING R

The code for an analysis in R is very parsimonious indeed:

```
summary(lm(mass ~ pop * length))
> summary(lm(mass ~ pop * length))
[...]
```

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-246.6623	24.8588	-9.923	5.72e-10	***
popMassif Central	197.5543	37.8767	5.216	2.41e-05	***
popJura	241.5563	39.1405	6.172	2.24e-06	***
length	5.9623	0.4323	13.792	6.66e-13	***
popMassif Central:length	-3.8041	0.6631	-5.737	6.52e-06	***
popJura:length	-4.7114	0.6595	-7.144	2.20e-07	***

```
[...]
```

Residual standard error: 10.01 on 24 degrees of freedom
Multiple R-squared: 0.9114, Adjusted R-squared: 0.8929
F-statistic: 49.37 on 5 and 24 DF, p-value: 7.48e-12

These coefficients can directly be compared with the beta vector since we simulated the data exactly in the default effects format of a linear model specified in R. The residual standard deviation is called residual standard error by R.

```
beta.vec
cat("And the residual SD was 10 \n")

> beta.vec
[1] -250 150 200 6 -3 -4
> cat("And the residual SD was 10 \n")
And the residual SD was 10
```

11.4 ANALYSIS USING WinBUGS (AND A CAUTIONARY TALE ABOUT THE IMPORTANCE OF COVARIATE STANDARDIZATION)

In WinBUGS, I find it much easier to fit the means parameterization of the model, i.e., to specify three separate linear regressions for each mountain range. The effects (i.e., differences of intercept or slopes with reference to the Pyrenees) are trivially easy to recover as derived parameters by just adding a few WinBUGS code lines. This allows for better comparison between input and output values.

```

# Define model
sink("lm.txt")
cat("
model {

# Priors
  for (i in 1:n.group){
    alpha[i] ~ dnorm(0, 0.001)      # Intercepts
    beta[i] ~ dnorm(0, 0.001)       # Slopes
  }
  sigma ~ dunif(0, 100)            # Residual standard deviation
  tau <- 1 / ( sigma * sigma)

# Likelihood
  for (i in 1:n) {
    mass[i] ~ dnorm(mu[i], tau)
    mu[i] <- alpha[pop[i]] + beta[pop[i]]* length[i]
  }

# Derived quantities
# Define effects relative to baseline level
  a.effe2 <- alpha[2] - alpha[1]    # Intercept Massif Central vs. Pyr.
  a.effe3 <- alpha[3] - alpha[1]    # Intercept Jura vs. Pyr.
  b.effe2 <- beta[2] - beta[1]      # Slope Massif Central vs. Pyr.
  b.effe3 <- beta[3] - beta[1]      # Slope Jura vs. Pyr.

# Custom tests
  test1 <- beta[3] - beta[2]        # Slope Jura vs. Massif Central
}
",fill=TRUE)
sink()

# Bundle data
win.data <- list(mass = as.numeric(mass), pop = as.numeric(pop), length = length,
n.group = max(as.numeric(pop)), n = n)

# Inits function
inits <- function(){ list(alpha = rnorm(n.group, 0, 2), beta = rnorm(n.groups, 1,
1), sigma = rlnorm(1))}

# Parameters to estimate
parameters <- c("alpha", "beta", "sigma", "a.effe2", "a.effe3", "b.effe2",
"b.effe3", "test1")

# MCMC settings
ni <- 1200
nb <- 200
nt <- 2
nc <- 3

```

```
# Start Markov chains
```

```
out <- bugs(win.data, inits, parameters, "lm.txt", n.thin=nt, n.chains=nc,
n.burnin=nb, n.iter=ni, debug = TRUE)
```

This is a simple model that converges rapidly. We inspect the results and compare them with the “truth” in the data-generating random process as well as with the inference from R ...

```
print(out, dig = 3)           # Bayesian analysis
beta.vec                     # Truth in the data-generating process
summary(lm(mass ~ pop * length)) # The ML solution again
```

... and are perplexed! WinBUGS claims that the Markov chains have converged (see Rhat values), but we get totally different estimates from what we should! Remember that `alpha[1]` and `beta[1]` in WinBUGS correspond to the intercept and the `length` main effect in the analysis in R and `a.ffe2`, `a.ffe3`, `b.ffe2`, `b.ffe3` to the remaining terms of the analysis in R.

```
> print(out, dig = 3) # Bayesian analysis
Inference for Bugs model at "lm.txt", fit using WinBUGS,
3 chains, each with 1200 iterations (first 200 discarded), n.thin = 2
n.sims = 1500 iterations saved
```

	mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
alpha[1]	-100.524	31.918	-162.910	-122.400	-100.600	-79.105	-38.406	1.002	1500
alpha[2]	-16.420	26.238	-65.813	-34.060	-16.670	1.595	34.926	1.000	1500
alpha[3]	-2.613	26.127	-51.847	-20.862	-2.920	15.893	47.693	1.001	1500
beta[1]	3.442	0.557	2.336	3.066	3.443	3.813	4.540	1.002	1500
beta[2]	1.586	0.464	0.669	1.272	1.588	1.896	2.467	1.000	1500
beta[3]	1.205	0.438	0.367	0.885	1.207	1.503	2.059	1.002	1500
sigma	15.841	3.089	10.669	13.650	15.525	17.637	22.610	1.001	1500
a.ffe2	84.103	38.764	6.313	57.752	84.620	110.300	160.082	1.002	1500
a.ffe3	97.911	41.127	16.015	70.735	97.855	126.200	175.105	1.001	1500
b.ffe2	-1.856	0.684	-3.178	-2.327	-1.856	-1.401	-0.479	1.002	1500
b.ffe3	-2.237	0.708	-3.574	-2.716	-2.253	-1.777	-0.820	1.001	1500
test1	-0.381	0.633	-1.608	-0.782	-0.396	0.032	0.866	1.001	1500
deviance	249.071	8.258	232.847	243.300	249.000	254.800	264.552	1.000	1500

```
[ ... ]

> beta.vec           # Truth in the data-generating process
[1] -250 150 200 6 -3 -4
```

So `lm()` is able to recover the right parameter values, up to sampling and estimation error, but WinBUGS is not. Why is this?

The problem turns out to reside in the lack of standardization of the covariate `length`. In WinBUGS, it is always advantageous to scale

covariates so that their extremes are not too far away from zero; otherwise, there may be nonconvergence or other problems. And the ugly thing here is that from looking at the convergence diagnostics (Rhat), we would never have guessed that there was a problem!

This example illustrates how useful it is to check the consistency of one's inference from WinBUGS with other sources, e.g., estimates from a simpler, but similar model run in WinBUGS or maximum likelihood estimates from another software. Know thy model! Alternatively, we could also have plotted the estimated regression lines into the observed data and would have seen easily that something was wrong.

As a quick check that lack of standardization was indeed the problem, we repeat both the maximum likelihood and the Bayesian analysis using a normalized version of the length covariate. This is simple; we just need to redefine the `length` covariate in the data list we pass to WinBUGS and can rerun the same code as before.

```
# Data passed to WinBUGS
win.data <- list(mass = as.numeric(mass), pop = as.numeric(pop), length =
  as.numeric(scale(length)), n.group = max(as.numeric(pop)), n = n)

# Start Markov chains
out <- bugs(win.data, inits, parameters, "lm.txt", n.thin=nt, n.chains=nc,
  n.burnin=nb, n.iter=ni, debug = FALSE)

...

# Inspect results
print(out, dig = 3)
> print(out, dig = 3)
[ ... ]
```

	mean	sd	2.5%	25%	50%	75%	97.5%	Rhat	n.eff
alpha[1]	97.796	3.406	90.679	95.820	97.905	100.100	104.252	1.004	1500
alpha[2]	74.952	3.504	67.864	72.690	75.030	77.292	81.580	1.004	1200
alpha[3]	66.535	3.653	59.009	64.060	66.660	68.995	73.265	1.007	350
beta[1]	41.243	3.158	34.829	39.150	41.215	43.332	47.316	1.005	870
beta[2]	14.486	3.763	7.318	11.947	14.570	16.920	22.112	1.002	1500
beta[3]	8.868	3.820	1.635	6.401	8.803	11.322	16.816	1.000	1500
sigma	10.587	1.678	7.937	9.380	10.370	11.510	14.516	1.005	630
a.effe2	-22.844	4.756	-31.771	-25.962	-22.815	-19.718	-13.498	1.003	1200
a.effe3	-31.262	5.000	-41.826	-34.470	-31.260	-28.050	-21.384	1.005	400
b.effe2	-26.757	4.802	-36.065	-30.032	-26.800	-23.620	-16.660	1.002	1500
b.effe3	-32.375	4.907	-41.445	-35.853	-32.350	-29.345	-22.162	1.001	1500
test1	-5.618	5.334	-16.024	-9.007	-5.836	-2.101	5.108	1.001	1500
deviance	225.218	4.558	218.700	221.900	224.500	227.800	235.500	1.009	270

```
[ ... ]
```


Compare with MLEs (R output slightly edited):

```
print(lm(mass ~ pop * as.numeric(scale(length)))$coefficients, dig = 4)
...
> print(lm(mass ~ pop * as.numeric(scale(length)))$coefficients, dig = 4)
      (Intercept)      popMassif Central
      98.94      -22.95
      popJura      as.numeric(scale(length))
     -31.54      41.78
popMassif Central:as.numeric(scale(length))  popJura:as.numeric(scale(length))
     -26.66      -33.01
```

Indeed, we now get consistent estimates in both analyses. Hence, previously the scale of the covariate didn't allow WinBUGS to converge, even though the Rhat values reported did indicate convergence. As a cautionary principle, we might therefore always consider to transform all covariates for WinBUGS, even if that slightly complicates presentation of results afterwards (for instance, in graphics). Transforming can mean centering, i.e., subtracting the mean, which changes the intercept only, but not the slope. Transforming can also mean normalizing, i.e., subtracting the mean and dividing the result by the standard deviation of the original covariate values. This changes both the intercept and the slope relative to an analysis with the original covariate. In the above case, centering will also work (want to try this out?).

11.5 SUMMARY

In a key chapter for your understanding of the modeling of grouped data, we have looked at the general linear model, or ANCOVA, in R and WinBUGS. We have focused on a model with one discrete and one continuous predictor and their interaction. Understanding ANCOVA is an important intermediate step to understanding the linear-mixed model in the next chapter.

EXERCISES

1. *Probability of a parameter*: What is the probability that the slope of the mass-length relationship of asp vipers is inferior in the Jura than in the Massif Central? Produce a graphical and a numerical answer.
2. *Related models*: Adapt the code to fit two variations of the model:
 - Fit different intercepts but a common slope
 - Fit the same intercept and the same slope

3. *Quadratic effects*: Add a quadratic term to the mass-length relationship, i.e., fit the model $\text{pop} + (\text{length} + \text{length}^2)$. You do not need to reassemble a data set that contains an effect of length squared, but you can simply take the data set we have already created in this chapter.
4. *Swiss hare data*: Fit an ANCOVA ($\text{pop} * \text{year}$, with year as a continuous explanatory variable) to the mean density. Also compute residuals and plot them to check for outliers.